

Software Requirements Specification (SRS)

Project: Invorto Voice AI Agent Platform (One-Phase GA Plan)

Version: 1.4\ **Date:** 21 Aug 2025 (IST)\ **Owner:** Invorto AI\ **Region:** AWS ap-south-1 (Mumbai)

0. Revision History

Version	Date	Author	Notes
1.4	2025-08-21	Invorto/ ChatGPT	Tool calling added (FR-TC1...TC10); webhook + timeline extended; schemaVersion bumped
1.3	2025-08-21	Invorto/ ChatGPT	Definitive vendor locks: Supabase Postgres (managed) replacing RDS option; diagrams & deploy steps updated; webhook schemaVersion bumped
1.2	2025-08-21	Invorto/ ChatGPT	Vapi parity fields; provider cost breakdown; assistant snapshot; destination; successEvaluation; webhook schemaVersion; cleanup dups
1.1	2025-08-21	Invorto/ ChatGPT	One-phase, definitive SRS for GA

1. Introduction

1.1 Purpose

Deliver a production-grade, low-latency, multi-tenant Voice AI platform in **one phase**. Includes telephony integration, real-time ASR→LLM→TTS agent runtime, tool calling, SDKs, webhooks, dashboard, cost guardrails, observability, and emotion-aware adaptation under feature flags.

1.2 Scope

- Inbound/outbound **SIP** via **Jambonz**; (optional) Browser WS client; WebRTC SFU deferred.
- Real-time pipeline: **Deepgram ASR** → **OpenAI GPT-4o-mini** → **Deepgram Aura-2 TTS** (defaults).
- **Native WebSockets** (binary PCM, 20–40 ms frames) for realtime.
- **Agent Runtime** with smart endpointing & barge-in, JSON-schema tools.
- **SDKs:** Node/TS server, Python, Browser WS realtime, Webhooks utils.
- **Data:** Redis (sessions/streams), Supabase Postgres (managed), S3 (recordings/transcripts/metrics).
- **Observability:** metrics, logs, traces, dashboards.
- **Compliance:** India DND/consent, PII redaction, ap-south-1 residency.
- **Emotion:** Energy + (optional) SER+NLP fusion; **feature-flagged**, non-blocking.

- **Costs:** All financial reporting & guardrails in **INR (₹)**.

1.3 Out of Scope (for GA)

- LiveKit/Pipecat/Daily
- WebRTC SFU dialer; Knowledge Base/RAG; self-hosted LLM/ASR/TTS; ClickHouse analytics (may be added post-GA).

1.4 Definitions & Abbreviations

- **ASR:** Automatic Speech Recognition.
- **TTS:** Text-to-Speech.
- **Barge-in:** Stopping TTS when user starts speaking.
- **EOU:** End-of-Utterance (turn boundary).
- **SER:** Speech Emotion Recognition.
- **SLO/SLA:** Service Level Objective/Agreement.
- **TTL:** Time To Live.
- **WS:** WebSocket.

2. Overall Description

2.1 Product Perspective

PSTN/SIP ↔ Jambonz (Media GW) ↔ Realtime WS Gateway ↔ ASR ↔ Agent (LLM + Tools + Policy) ↔ TTS ↔ Jambonz

transcripts/metrics)

└ Redis Streams (events)

└ Postgres (ops data)

└ S3 (recordings/

└ Webhooks → Tenant apps

└ Dashboard/SDKs

2.2 User Classes

- **Developers (Tenant):** Create agents, start calls, handle webhooks.
- **Supervisors/Analysts:** Review conversations, KPIs, spend, QA.
- **Ops/SRE:** Capacity, scaling, incidents, compliance.

2.3 Assumptions & Dependencies

- SIP trunks & numbers provisioned; AWS ap-south-1 available.
- Managed model providers reachable with low jitter; stable egress.
- Pricing/costing in ₹; exchange rates configurable.

2.4 Constraints

- Hard latency budgets (see NFR).
- Regulatory compliance (TRAI/DoT DND, recording disclosure).
- One-phase delivery: use **feature flags** for risky features.

2B. Bill of Materials (Definitive choices)

- **Cloud/Region:** AWS **ap-south-1** (Mumbai)
- **Media/Telephony:** **Jambonz** (self-hosted) with SIP trunks (Jio/Twilio, pluggable)
- **Realtime Gateway:** Custom WS service (Node/TS or Go) on **ECS Fargate**
- **ASR (streaming):** **Deepgram** Nova family
- **TTS (neural streaming):** **Deepgram Aura-2** (interruptible)
- **LLM (reasoning):** **OpenAI** `gpt-4o-mini` (tool use enabled)
- **State/Cache/Bus:** **Redis** (AWS ElastiCache) — keys/streams/locks
- **Relational DB:** **Supabase Postgres (managed)**
- **Object Store:** **Amazon S3** (encrypted, lifecycle rules)
- **IaC:** **Terraform**; **GitHub Actions** for CI/CD
- **Observability:** **OpenTelemetry** + CloudWatch; **Langfuse** for LLM tracing
- **Security:** AWS WAF, Secrets Manager, HMAC webhooks, JWT for WS

3. Functional Requirements (FR)

3.1 Telephony & Media IO

- **FR-T1** Inbound SIP via Jambonz; bridge RTP → WS PCM (8/16 kHz).
- **FR-T2** Outbound SIP; provider callbacks; call status updates.
- **FR-T3** DTMF detection (RFC2833/INFO) surfaced as events.
- **FR-T4** Audio frame policy: **20–40 ms**, mono PCM; single resample max.
- **FR-T5** Jitter buffer (40–80 ms), packet-loss concealment.
- **FR-T6** Consent prompt injection before free speech; configurable.

3.2 Realtime Gateway & Protocol

- **FR-R1** WebSocket endpoint `wss://api.../v1/realtime/:callId`; **JWT** or API key via header/`Sec-WebSocket-Protocol`.
- **FR-R2** Binary frames for audio; heartbeat ping/pong (15 s); sticky sessions behind ALB/Nginx.
- **FR-R3** Message types:

```
// client→server
{ "t":"start", "callId":"c_123", "agentId":"a_1", "locale":"en-IN" }
{ "t":"audio", "seq":123, "pcm16": "<binary>" }
{ "t":"tool.result", "id":"tk-1", "result": {"ok":true} }
// server→client
{ "t":"stt.partial", "text":"I would like...", "ts":170.4 }
{ "t":"stt.final", "text":"I want pricing", "ts":171.9 }
```

```
{ "t":"llm.delta", "text":"Sure, our plans..." }
{ "t":"tool.call", "id":"tk-1", "name":"book_meeting", "args":{"..."} }
{ "t":"tts.chunk", "seq":201, "pcm16":"<binary>" }
{ "t":"control.bargein", "action":"stop-tts" }
{ "t":"emotion.window", "energy_db":-16.3, "speaking":true }
{ "t":"emotion.state", "class":"neutral", "arousal":0.3, "valence":0.0,
"confidence":0.5 }
{ "t":"end", "reason":"hangup" }
```

3.3 ASR, LLM, TTS

- **FR-A1** ASR streaming (Deepgram): partials < 150 ms; finals on EOU.
- **FR-A2** LLM streaming (GPT-4o-mini): first tokens < 300 ms; function/tool call support.
- **FR-A3** TTS streaming (Aura-2): first audio < 500 ms; interruptible; pre-warm on call start.
- **FR-A4** Provider overrides per agent (ElevenLabs/Azure) via config; default Aura-2.

3.4 Turn-Taking & Endpointing

- **FR-E1** Smart endpointing provider toggle: `invorto` (default) | `livekit` | `off`.
- **FR-E2** Tunables: `silenceMs` (180–250), `minWords` (≥ 2), `confidenceThreshold` (0.6–0.7), `waitFunction` (e.g., `200 + 8000*x`).
- **FR-E3** Barge-in cutoff ≤ 150 ms; stop TTS immediately; preserve LLM stream.

3.5 Agent Runtime & Tools

- **FR-G1** Agent config CRUD: prompt, voice, locale, interruptible, temperature, guardrails, end_call rules.
- **FR-G2** Tools defined with **JSON Schema**; secure invocation with timeouts/retries/idempotency locks.
- **FR-G3** End-of-call summary + structured disposition (success, reason, next steps).
- **FR-G4** Policy engine (see Emotion): dynamic knobs for LLM/TTS/endpointing.

3.6 Webhooks & Integrations

- **FR-W1** Tenant-registered HTTPS endpoint; HMAC SHA-256 signatures.
- **FR-W2** Events: `call.started`, `asr.final`, `agent.tool_call`, `agent.summary`, `call.ended`, `billing.usage.updated`, `qa.flags`.
- **FR-W3** Retries (3× exponential); DLQ on failure.

3.7 SDKs & Specs

- **FR-S1** SDKs: `@invorto/server` (Node/TS), `invorto` (Python), `@invorto/realtime` (Browser WS), `@invorto/webhooks`.
- **FR-S2** Publish **OpenAPI (REST)** + **AsyncAPI (WS)**; versioned; SemVer.

3.8 Data & Storage

- **FR-D1** Redis for sessions (TTL), Streams for events, rate limits, locks.
- **FR-D2** Supabase Postgres (managed): tenants, agents, calls, event summaries, spend.

- **FR-D3** S3: recordings (WAV/MP3), transcripts (JSON), metrics NDJSON.
- **FR-D4** All costs & budgets in **INR (₹)**; per-tenant caps.
- **FR-D5 Per-provider cost breakdown** captured post-call: `usage.costs[]` with `{type, provider, minutes|units, costInr}`; top-level `usage.transportMinutes` and `usage.transportCostInr` included when telephony is used.

3.9 Emotion & QA (non-blocking)

- **FR-Q1** Energy meter (RMS→dBFS) every 250 ms; `emotion.window` events.
- **FR-Q2** (Flagged) SER + text sentiment; fused `emotion.state` (arousal, valence, class).
- **FR-Q3** Policy actions per label (frustrated/angry/confused/disengaged/positive) tuning LLM/TTS/endpointing for **next** turn.
- **FR-Q4** QA flags: `frustration_spike`, `toxicity`, `overlap_risk`.

4. Non-Functional Requirements (NFR)

4.1 Performance & Latency (p95 targets)

- **Call setup:** < 2.0 s after answer.
- **ASR first partial:** < 150 ms.
- **EOU detect:** 300–450 ms (short utterances).
- **LLM first token:** < 300 ms.
- **TTS first chunk:** < 500 ms.
- **User speak → bot speak:** ≤ 1.5 s total.
- **Barge-in cutoff:** ≤ 150 ms.

4.2 Scalability & Capacity

- 1,000 concurrent calls GA target; horizontal scale media + gateway; Redis in same AZ; async writers.
- Burst: 10k calls/hr for 5 min.

4.3 Reliability & Availability

- SLA 99.9% monthly (API + media).
- MTTR: Sev-2 < 1h; Sev-1 < 4h.
- Circuit breakers & provider failover.

4.4 Security & Privacy

- TLS 1.2+; JWT auth; HMAC webhooks; audit logs.
- PII redaction for transcripts; selective field masking.
- Secrets in AWS Secrets Manager; rotation 90d.
- RLS in Postgres; S3 SSE-KMS; India data residency by default.

4.5 Observability

- Metrics: `asr_first_partial_ms`, `eou_ms`, `llm_first_token_ms`, `tts_first_chunk_ms`, `bargein_cutoff_ms`, `cost_per_call_inr`.
- Logs: structured JSON with call/turn IDs.
- Traces: OpenTelemetry spans per turn.
- Dashboards: live Energy graph, calls, spend.

4.6 Data Retention & Residency

- Recordings default 180d; transcripts 365d; configurable.
- Right-to-erasure workflow; purge ≤ 14 days.
- ap-south-1 at rest; cross-region backups as per DR.

4.7 Accessibility & i18n

- en-IN/Hinglish baseline; Hindi support via ASR/LLM locales; TTS voice selection per agent.
 - Dashboard accessible (contrast, keyboard, ARIA labels).
-

5. System Architecture

5.1 Components

- **Jambonz Media GW:** SIP/RTP \leftrightarrow WS, VAD, DTMF.
- **Realtime WS Gateway:** binary WS, backpressure, heartbeats, auth.
- **ASR Adapter:** Deepgram streaming client.
- **Agent Runtime:** LLM orchestration, tool calling, endpointing, policy engine.
- **TTS Adapter:** Deepgram Aura-2 with pre-warm & caching.
- **Event Bus:** Redis Streams; sessions & rate limits in Redis.
- **Data Layer:** Supabase Postgres (managed), S3.
- **Webhooks Service:** signer/dispatcher with retries & DLQ.
- **Dashboard:** Supabase + Metabase (initial).
- **SDKs:** server, python, realtime, webhooks.

5.2 Deployment (AWS ap-south-1)

- **API/WS services:** ECS Fargate or EC2 ASG; ALB (WS upgrade).
- **Jambonz media:** EC2 with autoscale; NLB/ALB as needed.
- **Redis:** ElastiCache in same AZ; Multi-AZ later.
- **DB:** Supabase Postgres (managed); PITR enabled.
- **Storage:** S3 with versioning & lifecycle rules.
- **Network:** VPC with private subnets; NAT for egress; SG least privilege; WAF on ALB.

5.3 Configuration & Feature Flags

```
{
  "providers": { "asr": "deepgram", "llm": "openai-4o-mini", "tts": "aura-2" },
  "endpointing": { "provider": "invorto", "silenceMs": 220, "minWords": 2,
"confidence": 0.65, "waitFunction": "200 + 8000 * x" },
  "features": { "emotion": false, "providerOverride": true, "rag": false },
  "cost": { "currency": "INR", "rate_per_usd": 87 }
}
```

6. Data Model & Schemas

6.1 Postgres (core) — backed by Supabase Postgres

```
create table tenants(
  id uuid primary key, name text not null, created_at timestampz default now()
);
create table api_keys(
  id uuid primary key, tenant_id uuid references tenants, hash text not null,
  role text check (role in ('admin', 'dev', 'analyst')), active boolean default
true,
  created_at timestampz default now()
);
create table agents(
  id uuid primary key, tenant_id uuid references tenants, name text,
  version int default 1, config jsonb not null,
  created_at timestampz default now(), updated_at timestampz default now()
);
create table calls(
  id uuid primary key, tenant_id uuid references tenants, agent_id uuid
references agents,
  direction text, from_num text, to_num text, started_at timestampz, ended_at
timestampz,
  status text, duration_sec int, cost_inr numeric(12,2) default 0
);
create table events(
  id bigserial primary key, call_id uuid references calls, ts timestampz,
  kind text, payload jsonb
);
create materialized view call_stats as
select tenant_id, date_trunc('day', started_at) d, count(*) n_calls,
       sum(duration_sec) dur, avg(duration_sec) avg_dur, sum(cost_inr) cost_inr
from calls group by 1,2;
```

6.1.1 Billing Costs (per-provider)

```
create table call_costs(  
  id bigserial primary key,  
  call_id uuid references calls on delete cascade,  
  type text check (type in ('transport','stt','tts','llm','platform','other'))  
not null,  
  provider text,  
  minutes numeric(10,2), -- for transport  
  units numeric(12,2), -- generic units (e.g., tokens, chars)  
  cost_inr numeric(12,2) not null default 0,  
  created_at timestamptz default now()  
);  
create index call_costs_call_id_idx on call_costs(call_id);
```

6.2 Redis (keys/streams)

```
# Sessions (TTL)  
set call:{id}:session {json} EX 7200  
# Events  
XADD events:{callId} * kind "stt.final" payload "{...}"  
XGROUP CREATE events:{callId} agent 0 MKSTREAM  
# Emotion  
XADD emo:{callId} * v -0.35 a 0.78 c frustrated conf 0.76 energy_db -14.2  
# Rate limits  
INCRBY ratelimit:{key}:{bucket} 1; EXPIRE ...  
# Tool locks  
SETNX toollock:{callId}:book_meeting 1 EX 10
```

6.3 S3 Layout

```
recordings/{callId}.wav  
transcripts/{callId}.json  
metrics/{callId}.ndjson
```

7. APIs (Summary)

7.1 REST

- `POST /v1/agents` → create agent {id, version}
- `GET /v1/agents/:id` / `PATCH /v1/agents/:id`
- `POST /v1/calls` → start outbound {id,status}

- GET /v1/calls/:id / GET /v1/calls (filters)
- POST /v1/tools/:name/invoke (server-side)
- GET /v1/metrics (tenant summary)

7.2 WebSocket (see §3.2)

- Auth via JWT or API key; ping/pong; binary audio.
- Close idle > 2 min; error codes: invalid_api_key, quota_exceeded, provider_unavailable, bad_request, conflict.

7.3 Webhooks

- Header: X-Signature: sha256=<hmac>; retries ×3; DLQ topic/log.

7.4 Webhook Payloads (Schemas)

All payloads include type, schemaVersion, callId, tenantId, ISO8601 timestamps, and idempotency key eventId. HMAC is sent in the header, not the body.

****7.4.1 **** agent.summary

```
{
  "type": "agent.summary",
  "schemaVersion": "agent.summary.v1.4",
  "eventId": "evt_01H...",
  "callId": "c_7f2b9",
  "tenantId": "t_acme",
  "startedAt": "2025-08-21T11:42:10Z",
  "endedAt": "2025-08-21T11:43:12Z",
  "durationSec": 62,
  "carrier": {
    "trunk": "jio-primary",
    "sipFinal": "200 OK",
    "q850": 16,
    "ani": "+911234567890",
    "dnis": "+91226789xxxx"
  },
  "media": { "avgMos": 4.2, "pktLossPct": 0.3, "jitterMsP95": 18 },
  "kpis": { "asrFirstPartialMs": 122, "eouMs": 340, "llmFirstTokenMs": 210,
    "ttsFirstChunkMs": 380 },
  "usage": { "asrMin": 1.0, "ttsChars": 360, "llmIn": 420, "llmOut": 150,
    "costInr": 1.63 },
  "summary": {
    "outcome": "meeting_booked",
    "nextStep": "Calendar invite sent",
    "notes": "User prefers morning slot."
  },
}
```

```

    "emotion": { "enabled": true, "dominant": "positive", "valence": 0.42,
"arousal": 0.31 },
    "artifacts": {
      "recording": "https://signed.s3/.../recordings/c_7f2b9.wav",
      "transcript": "https://signed.s3/.../transcripts/c_7f2b9.jsonl",
      "metrics": "https://signed.s3/.../metrics/c_7f2b9.ndjson",
      "sipPcap": "https://signed.s3/.../pcap/c_7f2b9_sip.pcap",
      "rtpPcap": null,
      "bundleZip": "https://signed.s3/.../bundles/c_7f2b9_artifacts.zip?
expires=24h"
    }
  }
}

```

****7.4.2 **** call.started

```

{ "type": "call.started", "eventId": "evt_...", "callId": "c_...",
"tenantId": "t_...",
  "startedAt": "...", "direction": "inbound|outbound", "agentId": "a_...",
  "from": "+91...", "to": "+91...", "carrier": {"trunk": "...",
"sipCallId": "..."} }

```

****7.4.3 **** call.ended

```

{ "type": "call.ended", "eventId": "evt_...", "callId": "c_...", "endedAt": "...",
"reason": "hangup|error|transfer", "sipFinal": "486 Busy Here", "q850": 17,
"durationSec": 62, "media": {"avgMos": 4.2} }

```

****7.4.4 **** billing.usage.updated

```

{ "type": "billing.usage.updated", "tenantId": "t_...",
"periodStart": "2025-08-01",
  "periodEnd": "2025-08-31", "totals": {"asrMin": 1234.5, "ttsChars": 812345,
  "llmIn": 4 500 000, "llmOut": 1 200 000, "costInr": 8123.45}, "atRisk":
false }

```

****7.4.5 **** qa.flags

```

{ "type": "qa.flags", "callId": "c_...", "flags":
[{"kind": "frustration_spike", "ts": 170.2, "level": 0.86}] }

```

7.5 Pull APIs for Analytics & Artifacts

- GET /v1/calls/{callId} → structured analytics (carrier/media/usage/summary/emotion)
- GET /v1/calls/{callId}/artifacts → signed URLs: recording, transcript.jsonl, metrics.ndjson, sip.pcap, rtp.pcap?, bundle.zip
- GET /v1/calls/{callId}/timeline → ordered event list (stt.partial/final, llm.delta marks, tts.chunk stamps, barge-ins, tool calls)

7.6 Error Codes (expanded)

Code	HTTP	Meaning
invalid_api_key	401	Key missing/disabled
quota_exceeded	429	Tenant limits breached
provider_unavailable	503	Upstream ASR/LLM/TTS down
bad_request	400	Validation/schema error
conflict	409	State conflict
not_found	404	Call or artifact not found

7.7 Vapi Parity Additions (Fields & Examples)

The following fields are **added** to achieve near 1:1 parity with Vapi's post-call object while staying off the hot path.

7.7.1 Usage: transport & provider breakdown

```
"usage": {
  "transportMinutes": 1.0,
  "transportCostInr": 0.00,
  "asrMin": 1.0,
  "ttsChars": 360,
  "llmIn": 420,
  "llmOut": 150,
  "costInr": 1.63,
  "costs": [
    { "type": "transport", "provider": "twilio", "minutes": 1.0, "costInr": 0.00 },
    { "type": "stt", "provider": "deepgram", "units": 1.0, "costInr": 0.80 }
  ]
}
```

- Present only when telephony/transport is used; otherwise `transport*` values are `0`.

7.7.2 Messages echo (for convenience)

```
"artifact": {
  "messages": [
    { "role": "user", "message": "Hi, pricing?", "time": 0.42, "endTime": 3.1,
      "secondsFromStart": 0.42, "duration": 2.68, "isFiltered": false,
      "detectedThreats": [] },
    { "role": "assistant", "message": "Plans start at...", "time": 3.2,
      "endTime": 9.7, "secondsFromStart": 3.2, "duration": 6.5 }
  ]
}
```

This is an **echo** derived from the transcript for easy consumption. The full transcript remains available at `transcripts/{callId}.jsonl`.

7.7.3 Assistant snapshot at call start

```
"agentConfigAtCallStart": {
  "asr": { "provider": "deepgram", "language": "en-IN", "confidenceThreshold":
    0.65 },
  "llm": { "provider": "openai", "model": "gpt-4o-mini", "temperature": 0.6,
    "maxTokens": 256 },
  "tts": { "provider": "deepgram", "voiceId": "aura-2", "rate": 1.0 },
  "endpointing": { "provider": "invorto", "silenceMs": 220, "minWords": 2 }
}
```

7.7.4 Destination / transfer summary (if used)

```
"destination": {
  "type": "number",
  "mode": "blind-transfer",
  "timeout": 60,
  "sipVerb": "refer",
  "holdAudioUrl": "...",
  "transferCompleteAudioUrl": "...",
  "summaryPlan": { "enabled": true, "timeoutSeconds": 42 }
}
```

7.7.5 Optional success evaluation

```
"analysis": {
  "summary": "Booked a demo",
  "structuredData": { "bookingId": "bk_123" },
}
```

```

    "successEvaluation": { "rubric": "NumericScale", "score": 4 }
  }

```

All additions are emitted in the `agent.summary` webhook (and available via `GET /v1/calls/{id}`) without affecting latency.

3.10 Capture & Retention Policy (Config + FR)

- **FR-CP1** Capture policy is per-tenant: `sipPcap.onFailure`, `sipPcap.onSuccess`, `rtpPcap.enabled`, `rtpPcap.onLowMosThreshold`, `recording.stereo`, retention days.
- **FR-CP2** PCAP capture is **off** by default except SIP on failures; RTP PCAP may be auto-enabled when MOS < threshold.
- **FR-CP3** All links provided to tenants are **time-limited signed URLs**.

Config example

```

"capturePolicy": {
  "sipPcap": { "onFailure": true, "onSuccess": false, "retentionDays": 30 },
  "rtpPcap": { "enabled": false, "onLowMosThreshold": 3.2, "retentionDays": 7 },
  "recording": { "stereo": true, "retentionDays": 180 }
}

```

8. Emotion Policy & Adaptation

8.1 Policy Config

```

"emotionPolicy": {
  "thresholds": {
    "frustrated": { "arousal": 0.75, "valence": -0.30, "durationMs": 3000 },
    "angry":      { "arousal": 0.85, "valence": -0.40, "durationMs": 2000 },
    "disengaged": { "arousal": 0.30, "valence": 0.00, "silenceMs": 3000 }
  },
  "actions": {
    "frustrated": { "llm": { "temperature": 0.2, "style": "empathetic_concise" },
                  "tts": { "style": "empathetic", "rate": 0.9 },
                  "endpointing": { "silenceMsDelta": -40, "bargeIn":
"high" } },
    "confused":  { "llm": { "style": "clarify_steps" },
                  "endpointing": { "waitFunction": "250 + 9000*x" } },
    "angry":     { "llm": { "style": "deescalate", "max_tokens": 80 },
                  "endpointing": { "bargeIn": "max" } }
  }
}

```

8.2 Behavior

- Emotion runs in **parallel**; never blocks ASR/LLM/TTS.
 - Policy deltas applied to **next** turn only.
 - Hysteresis to avoid flip-flop (min dwell 1.2 s).
-

9. Costing & Guardrails (₹)

- Record **ASR mins**, **TTS chars**, **LLM tokens** per call; compute **₹ cost** with provider rates + fx.
 - Per-tenant caps (daily/weekly); 429 quota_exceeded on breach.
 - Alerts at 80%/100% of caps; webhook billing.usage.updated.
-

10. Quality, Testing & Acceptance

10.1 Test Types

- Unit (providers, tools, parsers).
- Integration (ASR/LLM/TTS loop; barge-in).
- E2E (Jambonz→summary).
- Load (≥ 50 cc soak; GA run @ target).
- Chaos (provider outages, packet loss 5–10%).

10.2 Acceptance Criteria (GA Gate)

- **Latency**: Speak start ≤ 1.5 s p95; barge-in ≤ 150 ms p95 on scripted set.
 - **Stability**: 2-hour soak @ ≥ 50 cc, zero fatal errors; memory/FD stable.
 - **Functionality**: streaming loop + tool call demo; webhooks HMAC + retries OK.
 - **Dashboard**: live Energy graph + call stats; spend in ₹.
 - **Security/Compliance**: JWT, HMAC, PII redaction; disclosure prompt; data residency ap-south-1.
 - **SDKs**: Quickstarts pass (Node, Python, Browser WS).
-

11. Operations

11.1 CI/CD

- GitHub Actions: build, unit tests, image scan, deploy staging→prod with approvals.
- Canary deploys; feature flags for risky features.

11.2 Monitoring & Alerts

- SLO alerts for p95s; provider error rates; Redis/DB saturation; spend guardrails.
- Synthetic call checks hourly.

11.3 Backup & DR

- Postgres PITR; S3 versioning; weekly cross-region backup.
- RTO 8 h / RPO 1 h (GA).

11.4 Runbooks

- Provider failover; Jambonz incident; Redis hot key; webhook backlog; barge-in misfires.
-

12. Security & Compliance

- OWASP ASVS L2; CIS hardening for hosts.
 - Access control: RBAC (admin/dev/analyst); audit logs for admin actions.
 - DND/consent workflows; immutable log of consent prompts.
-

13. Risk Register (Top)

Risk	Impact	Mitigation
ASR/TTS outage	Call failures	Multi-provider fallback; circuit breaker
TTS latency spikes	UX degrade	Pre-warm; cache canned phrases; switch provider
Cost spikes	Budget breach	Per-tenant ₹ caps; cheap TTS fallback
SIP trunk issue	No calls	Multi-trunk routing; health checks
PII leakage	Compliance	Redaction; access controls; logs scrubbed

14. Deliverables Checklist (One-Phase)

-

15. Appendices

A. Environment Variables (sample)

```
INVORTO_ENV=prod
REGION=ap-south-1
DB_URL=...
REDIS_URL=...
S3_BUCKET_RECORDINGS=...
```

```
S3_BUCKET_TRANSCRIPTS=...
ASR_PROVIDER=deepgram
LLM_PROVIDER=openai-4o-mini
TTS_PROVIDER=aura-2
WEBHOOK_SECRET=...
JWT_PUBLIC_KEY=...
CURRENCY=INR
USD_INR_RATE=87
```

B. Style Snippets (LLM)

```
[empathetic_concise] Acknowledge the issue in one short sentence. Offer one
actionable next step.
[clarify_steps] Confirm understanding in ≤10 words, then list at most 3 steps.
[deescalate] Apologize once, state what you will do now, avoid blame.
[cta_short] Ask permission and propose one clear next action.
```

C. Nginx WS Snippet

```
proxy_set_header Upgrade $http_upgrade;
proxy_set_header Connection "upgrade";
proxy_read_timeout 3600s;
proxy_buffering off;
```

5A. System Design Diagrams

5A.1 High-Level Architecture (Flow)

```
flowchart LR
    A[PSTN/SIP Caller] -->|RTP/SIP| B[Jambonz Media GW]
    subgraph AWS [ap-south-1]
        B -->|PCM 20-40ms (WS)| C[Realtime WS Gateway]
        C --> D[ASR Adapter (Deepgram)]
        D --> E[Agent Runtime  
(LLM + Tools + Policy)]
        E --> F[TTS Adapter (Aura-2)]
        F -->|PCM| B
        C <-->|events| G[(Redis Streams)]
        E -->|events, summaries| H[(Supabase Postgres)]
        C -->|recordings, transcripts, metrics| I[(S3)]
        C --> J[Webhooks Dispatcher]
        J --> K[Tenant Apps]
        L[Dashboard] --> H
    end
```



```
L --> G
end
```

5A.2 Deployment Topology

```
flowchart TB
    subgraph VPC
        subgraph Public Subnet
            ALB[ALB (TLS/WS)]
        end
        subgraph Private Subnets
            ECS1[ECS/Fargate: Realtime WS]
            ECS2[ECS/Fargate: Orchestrator/API]
            EC2J[EC2 ASG: Jambonz Media]
            REDIS[ElastiCache Redis]
            SUPA[(Supabase Postgres (managed))]
            S3[(S3 Buckets)]
        end
    end
    Internet --> ALB --> ECS1 --> REDIS
    ALB --> ECS2 --> SUPA
    PSTN --> EC2J --> ECS1
    ECS1 --> S3
    ECS2 --> S3
```

5A.3 Inbound Call Sequence (Detailed)

```
sequenceDiagram
    participant PSTN
    participant J as Jambonz
    participant RT as Realtime WS
    participant ASR as ASR (Deepgram)
    participant AG as Agent (LLM+Tools)
    participant TTS as TTS (Aura-2)
    participant R as Redis Streams
    participant DB as Postgres
    PSTN->>J: INVITE / RTP
    J->>RT: WS start(callId, agentId)
    RT->>ASR: streaming PCM frames (20-40ms)
    ASR-->>RT: stt.partial (<150ms)
    ASR-->>RT: stt.final (+EOU)
    RT->>AG: final transcript
    AG-->>RT: tool.call (if any)
    RT-->>R: XADD events: agent.tool_call
    Note over AG: stream LLM tokens
```

```

AG-->>RT: llm.delta (stream)
RT->>TTS: synth stream
TTS-->>RT: tts.chunk (first <500ms)
RT->>J: PCM out (interruptible)
RT->>R: XADD emotion.window/state (parallel)
RT->>DB: summary + cost (async writer)

```

5A.4 Emotion→Policy Adaptation

```

sequenceDiagram
    participant Emo as Emotion Worker
    participant R as Redis
    participant RT as Realtime WS
    participant AG as Agent Runtime
    Emo->>R: XADD emo:{callId} {arousal,valence,energy}
    R-->>RT: consume emo events (consumer group)
    RT->>AG: policy.update (adjust llm/tts/endpointing for next turn)

```

16. Gap Review & Updates (What we added/ensured)

Status: Core is complete. We added/clarified the following to make the SRS exhaustive:

1. **Call control** (added): outbound **DTMF send** and **transfer/hand-off** requirements.
2. **Provider failover policy** (clarified): circuit breakers + hot config switches for ASR/TTS/LLM.
3. **India compliance** (expanded): DPDP Act 2023 alignment; consent/opt-out ledger.
4. **Security** (expanded): optional **IP allowlists** per tenant; optional **SRTP/SIPS** where trunks support it; per-tenant KMS keys.
5. **Audio hygiene** (added): AGC/normalization, silence auto-hangup rules.
6. **Billing** (added): per-tenant **GST-ready invoices** in ₹; usage exports.
7. **Redis HA note** (clarified): Multi-AZ/auto-failover recommended at 1k cc.
8. **API throttles** (clarified): default per-key limits and error codes.

16.1 New/Updated Requirements

- **FR-T7** Outbound DTMF: send RFC2833/SIP INFO digits programmatically (useful for IVR navigation).
- **FR-T8** Call Transfer: support **blind** and **attended** transfer; optional conference bridge and whisper.
- **FR-SEC4** IP allowlists per tenant for REST/WS; toggled from console.
- **FR-SEC5** Optional **SRTP/SIPS** end-to-end where the trunk supports.
- **FR-C1** DPDP Act 2023 alignment: consent capture log, purpose limitation, erasure workflow.
- **FR-BILL1** Invoicing: monthly invoice in ₹ with **GST** fields; CSV usage export.
- **FR-AUDIO1** Audio normalization/AGC on TTS output and inbound mic where applicable.
- **FR-END1** Auto-hangup after configurable long silence/inactivity window (e.g., 45–90 s).
- **FR-RL1** API rate limits: default 60 req/min per key (REST), 3 WS connects/sec per key; 429 with `quota_exceeded`.

(All above inherit the same latency/observability constraints and are non-blocking to the hot audio path.)

Note:* If you want live-listen/whisper for supervisors, that's an optional Phase-later enhancement and remains out of GA scope.*

7.8 Tool Calling (In-call HTTP APIs)

Goal: Let the agent securely call external/internal APIs *during a live call* for lookups and actions, without breaking the ≤ 1.5 s p95 turn budget.

Functional Requirements

- **FR-TC1 (Tool definition):** Declarative tools with `{name, method, url, headers, inputSchema, outputSchema}`. Secrets via **AWS Secrets Manager**. Only **allow-listed** domains.
- **FR-TC2 (Blocking vs non-blocking):** Blocking tools have ≤ 800 ms latency budget on the hot path; non-blocking tools stream speech immediately and apply results next turn.
- **FR-TC3 (Timeouts & retries):** Default timeout **2 s**; exponential backoff $\times 2$; per-tool circuit breaker.
- **FR-TC4 (Idempotency & audit):** Per-turn idempotency key; every call logged with `{name, ok, latencyMs, status}` in **timeline** and summarized in **agent.summary.toolCalls[]**.
- **FR-TC5 (LLM mediation):** LLM may propose `tool(name, args)`; the orchestrator validates against schema/policy before execution.
- **FR-TC6 (Response limits & safety):** Max response 64 KB; JSON parsed; PII filtered per tenant policy; never pipe raw API text directly to TTS.
- **FR-TC7 (Fallback UX):** On error/timeout the agent uses a graceful utterance and optionally queues a non-blocking retry.
- **FR-TC8 (Security):** TLS only; optional outbound **HMAC**; per-tool rate limits; tenant-scoped secrets.
- **FR-TC9 (Latency safeguards):** If a blocking tool exceeds budget, it's cancelled and converted to non-blocking.
- **FR-TC10 (DevX):** `/v1/tools/validate` for schema checks; sandboxed test execution in staging.

Tool definition (example)

```
{
  "name": "get_order_status",
  "method": "GET",
  "url": "https://api.example.com/orders/{orderId}",
  "headers": {"Authorization": "Bearer ${secrets.crm_api}"},
  "inputSchema": {"type": "object", "properties": {"orderId": {"type":
"string"}}}, "required": ["orderId"]},
  "outputSchema": {"type": "object", "properties": {"status":
{"type": "string"}, "eta": {"type": "string"}}}
}
```

Webhooks & APIs

- **agent.summary (v1.4):** now includes `toolCalls[]` (see 7.4.6) and keeps `schemaVersion` updated.
- **/v1/calls/{id}/timeline:** adds `tool.requested`, `tool.succeeded`, `tool.failed` events.
- **/v1/tools/validate:** POST to validate a tool definition or a sample invocation.

Timeline event shapes

```
{ "ts":"2025-08-21T11:43:00Z", "type":"tool.requested",  
  "name":"get_order_status", "args":{"orderId":"A123"}, "blocking":true }  
{ "ts":"2025-08-21T11:43:00Z", "type":"tool.succeeded",  
  "name":"get_order_status", "latencyMs":312, "httpStatus":200,  
  "resultHash":"sha256:..." }  
{ "ts":"2025-08-21T11:43:00Z", "type":"tool.failed", "name":"get_order_status",  
  "latencyMs":1000, "error":"timeout" }
```

agent.summary excerpt (`toolCalls[]`)

```
"toolCalls": [  
  { "name":"get_order_status", "blocking": true, "ok": true, "latencyMs": 312,  
    "httpStatus": 200,  
    "startedAt":"2025-08-21T11:43:00Z", "endedAt":"2025-08-21T11:43:00Z",  
    "args": {"orderId":"A123"},  
    "resultSummary": {"status":"shipped", "eta":"2025-08-22"} }  
]
```

17. Deployment Plan (Definitive)

17.1 Environments & Release Cadence

- **Environments:** `dev` → `staging` → `prod` (all in **ap-south-1**).
- **Release model:** trunk-based; **semantic versioning**; canary in prod at **10%** traffic for 30 min → full rollout on green.
- **Change freeze:** 24h before major launches; emergency patch procedure allowed.

17.2 Infrastructure as Code (IaC)

- **Terraform** mono-repo with modules:
- `vpc` (3 AZs private; 2 AZs public), `nat`, route tables
- `ecs-cluster` (Fargate for **Realtime WS + Orchestrator/API**)
- `ec2-asg-jambonz` (Spot+On-Demand mix; UDP/RTP SG rules)
- `alb` (WS upgrade, WAF), `nlb` (optional for RTP)
- `redis` (ElastiCache)
- **Supabase (managed Postgres)** — provisioned outside AWS; connect via private peering or TLS; managed in a dedicated **Supabase project**

- `s3` (recordings/transcripts/metrics/bundles), lifecycle + SSE-KMS
- `secrets` (AWS Secrets Manager + SSM Parameter Store)
- `waf` (rate-limit, IP allowlists), `cloudwatch-alarms`
- **State:** Terraform Cloud or S3 + DynamoDB lock.

17.3 Build Artifacts

- **Containers:** `api`, `realtime`, `webhooks`, `workers`, `dashboard`, `jambonz` images → pushed to **ECR** with tag `vX.Y.Z+gitsha`.
- **Specs:** OpenAPI & AsyncAPI published to S3/Docs site.
- **Schemas:** `/schemas/agent.summary.v1.4.json` in artifact bundle.

17.4 CI/CD

- **GitHub Actions** workflows:
- **build-test** (lint, unit, integration, SBOM, image scan)
- **provision-dev** (Terraform plan/apply on merge to `main`)
- **deploy-staging** (ECS rolling update; DB migrations `--dry-run`)
- **smoke-staging** (synthetic call, webhook verification, SLO probes)
- **canary-prod** (ECS service w/ **weighted target groups** 10%)
- **promote-prod** (shift to 100% if KPIs < thresholds)
- **rollback** (redeploy previous TaskDef + revert config)

17.5 Provisioning Steps (per env)

1) Apply `vpc`, `alb`, `redis`, connect **Supabase Postgres**, and `s3`. 2) Create **KMS keys** (S3, DB, secrets) and rotate aliases. 3) Stand up **ECS services** (`realtime`, `api`, `webhooks`, `workers`, `dashboard`). 4) Create **EC2 ASG** for **jambonz** media nodes; attach NLB/SGs. 5) Issue TLS certs (ACM) → bind to ALB. 6) Configure **WAF** rules (rate-limit, IP allowlists). 7) Seed **Secrets Manager** entries (provider keys, HMAC secret, JWT keys).

17.6 Release/Runbook (prod)

- **Pre-deploy:** DB migrations `--plan`; feature flags default **OFF** (`emotion`, `providerOverride`, `rag`).
- **Deploy:** ECS rolling w/ `minHealthyPercent=100`, `maxPercent=200`.
- **Canary:** route 10% via weighted TG; watch **p95**: `user→bot ≤1.5s`, `barge-in ≤150ms`, error rate < 1% for 30 min.
- **Promote:** shift 100% traffic; tag release in Git.
- **Rollback:** single click to prior TaskDef; DB rollback via migration down or **feature-flag disable**.

17.7 Scaling & Capacity

- **Realtime WS (Fargate):** CPU 0.5 vCPU / 1GB; target 150 calls/task; scale on **active WS connections** & **CPU>60%**.
- **API/Webhooks/Workers:** 0.25–0.5 vCPU; scale on RPS & queue depth.
- **Jambonz ASG:** min 2, desired N; scale on RTP pps & CPU.
- **Redis:** cache.m6g.large single-AZ @ GA; Multi-AZ at 1k cc.

17.8 Networking & Security

- **Private subnets** for services; NAT only for egress.
- **SGs**: least-privilege (RTP/5060/WS/443 as needed).
- **JWT** for WS; **HMAC** for webhooks; **WAF** basic rules/ratelimits.
- **Audit logs** for admin actions; CloudTrail enabled.

17.9 Data & Migrations

- **Migrations**: Prisma/Alembic migrations run as init containers; `--dry-run` in staging, `--apply` in prod.
- **Backups**: Supabase PITR (WAL); S3 versioning; weekly cross-region copy.
- **Purge**: erasure jobs within 14 days; retention as per §4.6.

17.10 Observability & SLOs

- **OpenTelemetry** collector sidecar → vendor (or OTLP to Loki/Tempo).
- **CloudWatch** dashboards & alarms for `asr_first_partial_ms`, `euo_ms`, `llm_first_token_ms`, `tts_first_chunk_ms`, `bargein_cutoff_ms`, `cost_per_call_inr`.
- **Synthetic calls** hourly; alert on SLO breach.

17.11 Cost Controls (₹)

- Tagging: `CostCenter`, `Env`, `Service` on all resources.
- Budget alarms per env (₹); anomaly detection for >2× daily spend.

17.12 DR & Incident

- **RTO 8h / RPO 1h**; runbook for provider failover.
- **GameDays** quarterly: SIP trunk fail, Redis failover, packet loss.

17.13 Go/No-Go Checklist

- [] Terraform plan = clean; drift < 1%
- [] SLOs green in staging for 24h
- [] Synthetic call pass rate ≥ 99%
- [] Webhook HMAC verified by test app
- [] Cost guards (₹ caps) configured for tenants

End of One-Phase GA SRS v1.4